

Name	Type	Summary	Pros	Cons
kNN	C	The k nearest examples to the one that has to be predicted (nearest in Euclidean distance when plotted across its features) vote on which class the new example should fall in. The example gets 'democratically' assigned to the winning class.	<ul style="list-style-type: none"> • Simple • Makes no assumptions about underlying data distribution • Fast training phase 	<ul style="list-style-type: none"> • Does not produce a model, limiting the ability to understand how features are related to the class • Requires selection of an appropriate k • Slow classification phase • Nominal features and missing data require additional processing
Naïve Bayes	C	Applied Bayes theorem to the data to predict the probability of an output. Known as Naïve because of its naïve assumptions about the features being equally important and independent. Often used for text classification (spam filters).	<ul style="list-style-type: none"> • Simple and fast • Does well with noisy and missing data • Requires few examples for training, but still works well with large datasets • Easy to obtain the estimated probability for a prediction 	<ul style="list-style-type: none"> • Relies on the assumption of 'equally important and independent features' • Not ideal if there are many numeric features • Estimated probabilities are less reliable than the predicted classes
Decision Tree	C	Basically, a big flowchart with binary answers at each node up to the leaf node (result).	<ul style="list-style-type: none"> • Does well on most types of problems. • Does not require the user to specify the model in advance 	<ul style="list-style-type: none"> • Often biased toward splits on features having a large number of levels • It is easy to overfit or underfit the model

			<ul style="list-style-type: none"> • Excludes unimportant features • Can be used on both small and large datasets • Model that can be interpreted without a mathematical background 	<ul style="list-style-type: none"> • Can have trouble modelling some relationships due to reliance on axis-parallel splits • Small changes in the training data can result in large changes to decision logic • Large trees are hard to interpret
RIPPER Rule Learner	C	Works very similar to a decision tree, but makes rules out of all the possible paths taken from the root node to the output	<ul style="list-style-type: none"> • Generates easy to understand, human-readable rules • Efficient on large and noisy datasets • Generally, produces a simpler model than a comparable decision tree 	<ul style="list-style-type: none"> • May result in rules that seem to defy common sense or expert knowledge • Not ideal for working on numeric data • Might not perform as well as more complex models
Multiple Linear Regression	R	An equation in terms of the independent variable (features) that fits the training data as good as possible (trying not to overfit). The features of the future variable to predict are simply plugged into the equation (or plotted onto the regression line graph) to predict its dependent variable.	<ul style="list-style-type: none"> • Most common approach for modelling numeric data • Can be adapted to model almost any modelling task • Provides estimates of both the strength and size of the relationships among features and the outcome 	<ul style="list-style-type: none"> • Makes strong assumptions about data • The model's form must be specified by the user in advance • Does not handle missing data • Only works with numeric features, so categorical data requires extra processing • Requires some knowledge of

				statistics to understand the model
Regression Trees	R	Exactly like a decision tree, but to get a numerical prediction, they make predictions based on the average value of the examples that reach a leaf.	<ul style="list-style-type: none"> • Combines the strengths of decision trees with the ability to model numerical data • Does not require the user to specify the model in advance • Uses automatic feature selection, which allows the approach to be used with a very large number of features • May fit some types of data much better than linear regression • Does not require knowledge of statistics to interpret the model 	<ul style="list-style-type: none"> • Requires a large amount of training data • Difficult to determine the overall net effect of individual features on the outcome • Large trees can become more difficult to interpret than a regression model
Model Trees	R	Just like Regression Trees, but at each leaf they build a multiple linear regression model from the examples reaching that node.	<ul style="list-style-type: none"> • All the above • More powerful and more accurate predictions than regression trees 	<ul style="list-style-type: none"> • All the above • Much more complicated to interpret than regression trees

Neural Networks	C/R	Inputs (from nodes) are weighted per their importance (usually calculated with backpropagation) and summed into a new node. The sum is then fed into an activation function (sigmoid usually) that passes on the signal if activated (sum is > than a threshold). Multiple nodes can be layered to model more complex relationships.	<ul style="list-style-type: none"> • Capable of modelling more complex patterns than nearly any other algorithm • Makes few assumptions about the data's underlying relationships 	<ul style="list-style-type: none"> • Extremely computationally intensive and slow to train, particularly if the network topology is complex • Very prone to overfitting training data • Results impossible to interpret
Support Vector Machines (With non-linear kernel)	C/R	Uses flat hyperplanes to separate the data plotted in multi-dimensional space. The partitions created by the hyperplanes tend to be, more or less, homogeneous. Kernel functions are used to 'add' calculated features that were not present in the raw data. This helps redistribute the data in a larger dimensional space to find a better-fit flat hyperplane to partition it.	<ul style="list-style-type: none"> • Not too much affected by noisy data • Not prone to overfitting • Known to obtain very good results 	<ul style="list-style-type: none"> • Various Kernel functions must be tested to find the best (trial and error) • Very lengthy training for large datasets • Results impossible to interpret

Random Forest	C/R	<p>An ensemble of trees (500 by default in R). After bootstrap sampling the training data it selects random features to train each individual tree to increase diversity. The trees then 'vote' the prediction and one final prediction is made from these votes (average for numerical predictions). The votes can be weighted based on previous model performance.</p>	<ul style="list-style-type: none"> • Ensemble • Performs well on most problems. • Can handle noisy and missing data as well as categorical or continuous features. • 'Automatically' selects the most important features. • Can be used on data with extremely large number of features or examples. 	<ul style="list-style-type: none"> • Not easy to interpret. • Requires a lot of work to tune the model to the data.
----------------------	-----	--	---	---